

alzheimer-prediction

March 28, 2024

```
[3]: trainpath=r"D:/Winter Semester 3/AIML/Project/Alzheimer_image/Alzheimer_s_
      ↳Dataset/train"
      testpath=r"D:/Winter Semester 3/AIML/Project/Alzheimer_image/Alzheimer_s_
      ↳Dataset/test"
```

```
[4]: import tensorflow
      from tensorflow.keras.layers import Dense, Flatten, Input, Dropout
      from tensorflow.keras.models import Model
      from tensorflow.keras.preprocessing import image
      from tensorflow.keras.preprocessing.image import ImageDataGenerator as IDG,
      ↳load_img
      from tensorflow.keras.applications import Xception
      from tensorflow.keras import layers, models
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import
      ↳SeparableConv2D, BatchNormalization, GlobalAveragePooling2D
      from tensorflow.keras.models import load_model
      from PIL import Image
      from tensorflow.keras.applications.xception import preprocess_input
      import numpy as np
```

```
[5]: img_size=180
      img_size_dim=[180,180]
      dimension=(img_size,img_size)
      zoom=[0.99,1.01]
      bright=[0.8,1.2]
      fill_mode='constant'
      data_format='channels_last'
      dir=trainpath
      data_gen=IDG(rescale=1./
      ↳255,brightness_range=bright,zoom_range=zoom,data_format=data_format,fill_mode=fill_mode,hor
      train_data_gen=data_gen.
      ↳flow_from_directory(directory=dir,target_size=dimension,batch_size=6000,shuffle=False)
```

Found 5121 images belonging to 4 classes.

```
[6]: train_data,train_labels=next(train_data_gen)
      print(train_data.shape,train_labels.shape)
```

(5121, 180, 180, 3) (5121, 4)

```
[7]: #handling imbalanced classes
      from imblearn.over_sampling import SMOTE
      sm=SMOTE(random_state=47)
      train_data,train_labels=sm.fit_resample(train_data.
        ↪reshape(-1,img_size*img_size*3),train_labels)
      train_data=train_data.reshape(-1,img_size,img_size,3)
      print(train_data.shape,train_labels.shape)
```

(10240, 180, 180, 3) (10240, 4)

```
[8]: from sklearn.model_selection import train_test_split
      train_data,test_data,train_labels,test_labels = ␣
        ↪train_test_split(train_data,train_labels,test_size=0.2,random_state=47)
      train_data,val_data,train_labels,val_labels = ␣
        ↪train_test_split(train_data,train_labels,test_size=0.2,random_state=47)
```

```
[36]: import tensorflow as tf

      mt = [tf.keras.metrics.CategoricalAccuracy(name='Accuracy'), tf.keras.metrics.
        ↪AUC(name='AUC')]
```

```
[10]: import tensorflow as tf
      from tensorflow.keras import layers, models

      # Create a Sequential model
      model = models.Sequential()

      # Add the first convolutional layer with 32 filters, a 3x3 kernel, and ReLU␣
        ↪activation
      model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(180, 180,␣
        ↪3)))

      # Add a max pooling layer with a 2x2 pool size
      model.add(layers.MaxPooling2D((2, 2)))

      # Add another convolutional layer with 64 filters, a 3x3 kernel, and ReLU␣
        ↪activation
      model.add(layers.Conv2D(64, (3, 3), activation='relu'))

      # Add another max pooling layer with a 2x2 pool size
      model.add(layers.MaxPooling2D((2, 2)))
```

```

# Add a flattening layer to convert the 3D feature maps to 1D feature vectors
model.add(layers.Flatten())

# Add a fully connected layer with 64 units and ReLU activation
model.add(layers.Dense(64, activation='relu'))

# Add the output layer with 10 units (assuming 10 classes) and softmax
  ↪activation for classification
model.add(layers.Dense(4, activation='softmax'))

# Compile the model with the Adam optimizer and categorical cross-entropy loss
  ↪function
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=mt)

# Print a summary of the model architecture
model.summary()

```

c:\Users\harsa\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
flatten (Flatten)	(None, 118336)	0
dense (Dense)	(None, 64)	7,573,568
dense_1 (Dense)	(None, 4)	260

Total params: 7,593,220 (28.97 MB)

Trainable params: 7,593,220 (28.97 MB)

Non-trainable params: 0 (0.00 B)

```
[11]: history=model.  
      ↪fit(train_data,train_labels,validation_data=(val_data,val_labels),epochs=5)
```

```
Epoch 1/5  
205/205          75s 355ms/step -  
AUC: 0.7029 - Accuracy: 0.4559 - loss: 1.4442 - val_AUC: 0.9587 - val_Accuracy:  
0.7993 - val_loss: 0.4942  
Epoch 2/5  
205/205          76s 371ms/step -  
AUC: 0.9729 - Accuracy: 0.8410 - loss: 0.3944 - val_AUC: 0.9720 - val_Accuracy:  
0.8383 - val_loss: 0.3880  
Epoch 3/5  
205/205          77s 376ms/step -  
AUC: 0.9939 - Accuracy: 0.9324 - loss: 0.1886 - val_AUC: 0.9798 - val_Accuracy:  
0.8707 - val_loss: 0.3421  
Epoch 4/5  
205/205          72s 353ms/step -  
AUC: 0.9955 - Accuracy: 0.9380 - loss: 0.1554 - val_AUC: 0.9887 - val_Accuracy:  
0.9109 - val_loss: 0.2415  
Epoch 5/5  
205/205          73s 356ms/step -  
AUC: 0.9993 - Accuracy: 0.9763 - loss: 0.0629 - val_AUC: 0.9890 - val_Accuracy:  
0.9329 - val_loss: 0.2147
```

```
[37]: from keras.models import Model  
  
def freeze_layers(model):  
    for i in model.layers:  
        i.trainable = False  
        if isinstance(i, Model):  
            freeze_layers(i)  
    return model
```

```
[38]: model_frozen = freeze_layers(model)  
      #model_frozen.save('alzheimer_3.h5')
```

```
[39]: model_frozen.summary()
```

Model: "sequential"

Layer (type)

Output Shape

Param #

conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
flatten (Flatten)	(None, 118336)	0
dense (Dense)	(None, 64)	7,573,568
dense_1 (Dense)	(None, 4)	260

Total params: 22,779,662 (86.90 MB)

Trainable params: 0 (0.00 B)

Non-trainable params: 7,593,220 (28.97 MB)

Optimizer params: 15,186,442 (57.93 MB)

```
[15]: test_model=load_model(r"D:\Winter Semester 3\AIML\Project\alzheimer_3.h5")
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

WARNING:absl>Error in loading the saved optimizer state. As a result, your model is starting with a freshly initialized optimizer.

```
[16]: img=image.load_img(r"D:\Winter Semester 3\STS\Cat-2\hello.jpg")
img=img.resize((180,180))
```

```
[17]: x=image.img_to_array(img)
x=np.expand_dims(x,axis=0)
img_data=preprocess_input(x)
```

```
[18]: print(np.argmax(test_model.predict(x),axis=1)[0])
```

1/1 0s 113ms/step

2

```
[ ]: "D:\Winter Semester 3\AIML\Project\Alzheimer_image\Alzheimer_s_
↳Dataset\test\VeryMildDemented\26 (44).jpg"
"D:\Winter Semester 3\AIML\Project\Alzheimer_image\Alzheimer_s_
↳Dataset\test\NonDemented\26 (62).jpg"
```

```
[47]: import tensorflow as tf
from tensorflow.keras.preprocessing import image
import numpy as np

# Load the saved model
loaded_model = tf.keras.models.load_model(r"D:\Winter Semester_
↳3\AIML\Project\alzheimer_3.h5")

# Check the model's summary to understand the architecture and input shape
loaded_model.summary()

# Load and preprocess the new image
img_path = r"D:\Winter Semester 3\STS\Cat-2\hello.jpg" # Replace_
↳'path_to_new_image.jpg' with the actual path to your image
img = image.load_img(img_path, target_size=(224, 224)) # Resize the image to_
↳match input size
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
img_array /= 255. # Normalize pixel values
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

WARNING:absl>Error in loading the saved optimizer state. As a result, your model is starting with a freshly initialized optimizer.

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_21 (Conv2D)	(None, 87, 87, 64)	18,496
max_pooling2d_9 (MaxPooling2D)	(None, 43, 43, 64)	0
flatten_6 (Flatten)	(None, 118336)	0
dense_22 (Dense)	(None, 64)	7,573,568

dense_23 (Dense)

(None, 4)

260

Total params: 7,593,222 (28.97 MB)


Trainable params: 7,593,220 (28.97 MB)

Non-trainable params: 0 (0.00 B)


Optimizer params: 2 (12.00 B)

```
[51]: import matplotlib.pyplot as plt
from matplotlib.pyplot import figure

fig = plt.figure(figsize=(12, 6), dpi=80)
plt.plot(history.history['loss'], color='teal', label='loss')
plt.plot(history.history['val_loss'], color='orange', label='val_loss')
fig.suptitle('Loss', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```



```
[57]: fig = plt.figure(figsize=(12, 6), dpi=80)
plt.plot(history.history['Accuracy'], color='teal', label='accuracy')
plt.plot(history.history['val_Accuracy'], color='orange', label='val_accuracy')
fig.suptitle('Accuracy', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```



```
[65]: # Evaluate the model on test data
test_loss_accuracy = model.evaluate(test_data, test_labels)
y_pred = model.predict(test_data).argmax(axis=1)
predicted_class = np.argmax(y_pred)
test_loss = test_loss_accuracy[0]
test_accuracy = test_loss_accuracy[1]

print("Test Loss:", test_loss)
print("Test Accuracy:", test_accuracy)
```

```
64/64          6s 88ms/step - AUC:
0.9921 - Accuracy: 0.9377 - loss: 0.1867
64/64          5s 80ms/step
Test Loss: 0.21458502113819122
Test Accuracy: 0.9897463917732239
```

```
[61]: from sklearn import metrics
import pandas as pd
df = pd.DataFrame(
```



```

data = metrics.confusion_matrix(np.argmax(test_labels, axis=1), y_pred),
columns = ['0', '1', '2', '3'],
index = ['0', '1', '2', '3']
)
df

```

```

[61]:
      0    1    2    3
0  505    0    2    3
1    0  503    0    0
2   22    0  431   46
3   23    0   46  467

```

```

[68]: import seaborn as sns


sns.heatmap(df, cmap="Blues", annot=True, fmt='.0f')

```

```

[68]: <AxesSubplot:>

```



```

[82]: from sklearn.metrics import classification_report

# Assuming you have your test labels and predicted classes
# Replace test_labels and predicted_classes with your actual data

```

```

# Convert one-hot encoded test labels to integer format
test_labels_int = np.argmax(test_labels, axis=1)

# Generate classification report
report = classification_report(test_labels_int, predicted_classes,
    ↳target_names=class_names)

print("Classification Report:")
print(report)

```

Classification Report:

	precision	recall	f1-score	support
MildDemented	0.96	0.98	0.97	510
ModerateDemented	1.00	1.00	1.00	503
NonDemented	0.91	0.91	0.91	499
VeryMildDemented	0.92	0.90	0.91	536
accuracy			0.95	2048
macro avg	0.95	0.95	0.95	2048
weighted avg	0.95	0.95	0.95	2048