

isease-classification-sequential-2

March 28, 2024

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import tensorflow as tf
import matplotlib.pyplot as plt

import os
from distutils.dir_util import copy_tree, remove_tree

from PIL import Image
from random import randint

from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.metrics import matthews_corrcoef as MCC
from sklearn.metrics import balanced_accuracy_score as BAS
from sklearn.metrics import classification_report, confusion_matrix

import tensorflow_addons as tfa
from keras.utils.vis_utils import plot_model
from tensorflow.keras import Sequential, Input
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.layers import Conv2D, Flatten
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.applications.inception_v3 import InceptionV3
from tensorflow.keras.preprocessing.image import ImageDataGenerator as IDG
from tensorflow.keras.layers import SeparableConv2D, BatchNormalization,
    ↪MaxPool2D

print("TensorFlow Version:", tf.__version__)
```

TensorFlow Version: 2.4.1

```
[2]: base_dir = "/kaggle/input/alzheimers-dataset-4-class-of-images/Alzheimer_s_
    ↪Dataset/"
```

```

root_dir = "./"
test_dir = base_dir + "test/"
train_dir = base_dir + "train/"
work_dir = root_dir + "dataset/"

if os.path.exists(work_dir):
    remove_tree(work_dir)

os.mkdir(work_dir)
copy_tree(train_dir, work_dir)
copy_tree(test_dir, work_dir)
print("Working Directory Contents:", os.listdir(work_dir))

```

Working Directory Contents: ['MildDemented', 'NonDemented', 'ModerateDemented', 'VeryMildDemented']

```

[3]: WORK_DIR = './dataset/'

CLASSES = [ 'NonDemented',
            'VeryMildDemented',
            'MildDemented',
            'ModerateDemented']

IMG_SIZE = 176
IMAGE_SIZE = [176, 176]
DIM = (IMG_SIZE, IMG_SIZE)

```

```

[4]: #Performing Image Augmentation to have more data samples

ZOOM = [.99, 1.01]
BRIGHT_RANGE = [0.8, 1.2]
HORZ_FLIP = True
FILL_MODE = "constant"
DATA_FORMAT = "channels_last"

work_dr = IDG(rescale = 1./255, brightness_range=BRIGHT_RANGE, zoom_range=ZOOM,
↳data_format=DATA_FORMAT, fill_mode=FILL_MODE, horizontal_flip=HORZ_FLIP)

train_data_gen = work_dr.flow_from_directory(directory=WORK_DIR,
↳target_size=DIM, batch_size=6500, shuffle=False)

```

Found 6400 images belonging to 4 classes.

```

[5]: def show_images(generator, y_pred=None):
    """
    Input: An image generator, predicted labels (optional)
    """

```

Output: Displays a grid of 9 images with labels

```
"""


# get image labels
labels =dict(zip([0,1,2,3], CLASSES))

# get a batch of images
x,y = generator.next()

# display a grid of 9 images
plt.figure(figsize=(10, 10))
if y_pred is None:
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        idx = randint(0, 6400)
        plt.imshow(x[idx])
        plt.axis("off")
        plt.title("Class:{}".format(labels[np.argmax(y[idx])]))

else:
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(x[i])
        plt.axis("off")
        plt.title("Actual:{} \nPredicted:{}".format(labels[np.
↪argmax(y[i])],labels[y_pred[i]]))

# Display Train Images
show_images(train_data_gen)
```



```
[6]: #Retrieving the data from the ImageDataGenerator iterator
train_data, train_labels = train_data_gen.next()
```

```
[7]: #Getting to know the dimensions of our dataset
print(train_data.shape, train_labels.shape)
```

```
(6400, 176, 176, 3) (6400, 4)
```

```
[8]: #Performing over-sampling of the data, since the classes are imbalanced
sm = SMOTE(random_state=42)
```

```

train_data, train_labels = sm.fit_resample(train_data.reshape(-1, IMG_SIZE *
↳IMG_SIZE * 3), train_labels)

train_data = train_data.reshape(-1, IMG_SIZE, IMG_SIZE, 3)

print(train_data.shape, train_labels.shape)

```

(12800, 176, 176, 3) (12800, 4)

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py:72:
FutureWarning: Pass classes=[0 1 2 3] as keyword args. From version 1.0
(renaming of 0.25) passing these as positional arguments will result in an error
"will result in an error", FutureWarning)

[9]: *#Splitting the data into train, test, and validation sets*

```

train_data, test_data, train_labels, test_labels = train_test_split(train_data,
↳train_labels, test_size = 0.2, random_state=42)
train_data, val_data, train_labels, val_labels = train_test_split(train_data,
↳train_labels, test_size = 0.2, random_state=42)

```

```

[10]: def conv_block(filters, act='relu'):
        """Defining a Convolutional NN block for a Sequential CNN model. """

        block = Sequential()
        block.add(Conv2D(filters, 3, activation=act, padding='same'))
        block.add(Conv2D(filters, 3, activation=act, padding='same'))
        block.add(BatchNormalization())
        block.add(MaxPool2D())

        return block

```

```

[11]: def dense_block(units, dropout_rate, act='relu'):
        """Defining a Dense NN block for a Sequential CNN model. """

        block = Sequential()
        block.add(Dense(units, activation=act))
        block.add(BatchNormalization())
        block.add(Dropout(dropout_rate))

        return block

```

```

[12]: def construct_model(act='relu'):
        """Constructing a Sequential CNN architecture for performing the_
↳classification task. """

```

```

model = Sequential([
    Input(shape=(*IMAGE_SIZE, 3)),
    Conv2D(16, 3, activation=act, padding='same'),
    Conv2D(16, 3, activation=act, padding='same'),
    MaxPool2D(),
    conv_block(32),
    conv_block(64),
    conv_block(128),
    Dropout(0.2),
    conv_block(256),
    Dropout(0.2),
    Flatten(),
    dense_block(512, 0.7),
    dense_block(128, 0.5),
    dense_block(64, 0.3),
    Dense(4, activation='softmax')
], name = "cnn_model")

return model

```

[13]: *#Defining a custom callback function to stop training our model when accuracy goes above 99%*

```

class MyCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        if logs.get('val_acc') > 0.99:
            print("\nReached accuracy threshold! Terminating training.")
            self.model.stop_training = True

my_callback = MyCallback()

#EarlyStopping callback to make sure model is always learning
early_stopping = EarlyStopping(monitor='val_loss', patience=2)

```

[14]: *#Defining other parameters for our CNN model*

```

model = construct_model()

METRICS = [tf.keras.metrics.CategoricalAccuracy(name='acc'),
            tf.keras.metrics.AUC(name='auc'),
            tfa.metrics.F1Score(num_classes=4)]

CALLBACKS = [my_callback]

model.compile(optimizer='adam',
              loss=tf.losses.CategoricalCrossentropy(),
              metrics=METRICS)

```

```
model.summary()
```

Model: "cnn_model"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 176, 176, 16)	448
conv2d_1 (Conv2D)	(None, 176, 176, 16)	2320
max_pooling2d (MaxPooling2D)	(None, 88, 88, 16)	0
sequential (Sequential)	(None, 44, 44, 32)	14016
sequential_1 (Sequential)	(None, 22, 22, 64)	55680
sequential_2 (Sequential)	(None, 11, 11, 128)	221952
dropout (Dropout)	(None, 11, 11, 128)	0
sequential_3 (Sequential)	(None, 5, 5, 256)	886272
dropout_1 (Dropout)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
sequential_4 (Sequential)	(None, 512)	3279360
sequential_5 (Sequential)	(None, 128)	66176
sequential_6 (Sequential)	(None, 64)	8512
dense_3 (Dense)	(None, 4)	260

Total params: 4,534,996
Trainable params: 4,532,628
Non-trainable params: 2,368

```
[15]: #Fit the training data to the model and validate it using the validation data
      EPOCHS = 100

      history = model.fit(train_data, train_labels, validation_data=(val_data,
      ↪ val_labels), callbacks=CALLBACKS, epochs=EPOCHS)
```

Epoch 1/100

256/256 [=====] - 23s 55ms/step - loss: 1.8004 - acc:

0.2737 - auc: 0.5321 - f1_score: 0.2737 - val_loss: 1.6713 - val_acc: 0.2471 - val_auc: 0.4970 - val_f1_score: 0.1213

Epoch 2/100

256/256 [=====] - 13s 49ms/step - loss: 1.4250 - acc: 0.3321 - auc: 0.5971 - f1_score: 0.3307 - val_loss: 2.7988 - val_acc: 0.2568 - val_auc: 0.5768 - val_f1_score: 0.1027

Epoch 3/100

256/256 [=====] - 13s 49ms/step - loss: 1.0525 - acc: 0.5263 - auc: 0.7979 - f1_score: 0.5210 - val_loss: 1.6007 - val_acc: 0.3398 - val_auc: 0.6151 - val_f1_score: 0.2682

Epoch 4/100

256/256 [=====] - 13s 50ms/step - loss: 0.8075 - acc: 0.6314 - auc: 0.8812 - f1_score: 0.6164 - val_loss: 0.8203 - val_acc: 0.6265 - val_auc: 0.8751 - val_f1_score: 0.6104

Epoch 5/100

256/256 [=====] - 13s 50ms/step - loss: 0.7167 - acc: 0.6691 - auc: 0.9052 - f1_score: 0.6555 - val_loss: 2.2500 - val_acc: 0.3145 - val_auc: 0.5723 - val_f1_score: 0.2327

Epoch 6/100

256/256 [=====] - 13s 51ms/step - loss: 0.6619 - acc: 0.7046 - auc: 0.9192 - f1_score: 0.7015 - val_loss: 0.5792 - val_acc: 0.7500 - val_auc: 0.9401 - val_f1_score: 0.7422

Epoch 7/100

256/256 [=====] - 13s 51ms/step - loss: 0.5958 - acc: 0.7363 - auc: 0.9350 - f1_score: 0.7324 - val_loss: 1.0391 - val_acc: 0.5815 - val_auc: 0.8315 - val_f1_score: 0.5865

Epoch 8/100

256/256 [=====] - 13s 52ms/step - loss: 0.5959 - acc: 0.7432 - auc: 0.9368 - f1_score: 0.7426 - val_loss: 1.0476 - val_acc: 0.5464 - val_auc: 0.8276 - val_f1_score: 0.5357

Epoch 9/100

256/256 [=====] - 13s 52ms/step - loss: 0.5397 - acc: 0.7667 - auc: 0.9471 - f1_score: 0.7654 - val_loss: 0.4883 - val_acc: 0.7812 - val_auc: 0.9561 - val_f1_score: 0.7794

Epoch 10/100

256/256 [=====] - 13s 52ms/step - loss: 0.4880 - acc: 0.7864 - auc: 0.9559 - f1_score: 0.7867 - val_loss: 1.6974 - val_acc: 0.5093 - val_auc: 0.7485 - val_f1_score: 0.4413

Epoch 11/100

256/256 [=====] - 13s 51ms/step - loss: 0.4846 - acc: 0.7894 - auc: 0.9567 - f1_score: 0.7872 - val_loss: 1.1824 - val_acc: 0.6421 - val_auc: 0.8877 - val_f1_score: 0.6106

Epoch 12/100

256/256 [=====] - 13s 51ms/step - loss: 0.4496 - acc: 0.8060 - auc: 0.9625 - f1_score: 0.8044 - val_loss: 0.6092 - val_acc: 0.7666 - val_auc: 0.9472 - val_f1_score: 0.7510

Epoch 13/100

256/256 [=====] - 13s 51ms/step - loss: 0.4040 - acc:

0.8319 - auc: 0.9707 - f1_score: 0.8310 - val_loss: 0.4064 - val_acc: 0.8047 -
val_auc: 0.9689 - val_f1_score: 0.7931
Epoch 14/100
256/256 [=====] - 13s 51ms/step - loss: 0.3385 - acc:
0.8616 - auc: 0.9787 - f1_score: 0.8591 - val_loss: 0.5648 - val_acc: 0.7700 -
val_auc: 0.9452 - val_f1_score: 0.7704
Epoch 15/100
256/256 [=====] - 13s 51ms/step - loss: 0.3226 - acc:
0.8676 - auc: 0.9806 - f1_score: 0.8663 - val_loss: 1.9555 - val_acc: 0.5234 -
val_auc: 0.8089 - val_f1_score: 0.4310
Epoch 16/100
256/256 [=====] - 13s 51ms/step - loss: 0.3050 - acc:
0.8781 - auc: 0.9823 - f1_score: 0.8769 - val_loss: 0.8674 - val_acc: 0.6982 -
val_auc: 0.9293 - val_f1_score: 0.6433
Epoch 17/100
256/256 [=====] - 13s 51ms/step - loss: 0.2490 - acc:
0.9027 - auc: 0.9881 - f1_score: 0.9027 - val_loss: 0.3560 - val_acc: 0.8633 -
val_auc: 0.9820 - val_f1_score: 0.8584
Epoch 18/100
256/256 [=====] - 13s 51ms/step - loss: 0.2023 - acc:
0.9263 - auc: 0.9917 - f1_score: 0.9254 - val_loss: 2.3461 - val_acc: 0.3975 -
val_auc: 0.7062 - val_f1_score: 0.3268
Epoch 19/100
256/256 [=====] - 13s 51ms/step - loss: 0.1911 - acc:
0.9330 - auc: 0.9926 - f1_score: 0.9332 - val_loss: 0.3677 - val_acc: 0.8608 -
val_auc: 0.9786 - val_f1_score: 0.8584
Epoch 20/100
256/256 [=====] - 13s 51ms/step - loss: 0.1843 - acc:
0.9342 - auc: 0.9934 - f1_score: 0.9334 - val_loss: 0.2853 - val_acc: 0.8979 -
val_auc: 0.9868 - val_f1_score: 0.8961
Epoch 21/100
256/256 [=====] - 13s 51ms/step - loss: 0.1652 - acc:
0.9411 - auc: 0.9942 - f1_score: 0.9410 - val_loss: 0.3495 - val_acc: 0.8828 -
val_auc: 0.9818 - val_f1_score: 0.8833
Epoch 22/100
256/256 [=====] - 13s 51ms/step - loss: 0.1322 - acc:
0.9534 - auc: 0.9961 - f1_score: 0.9534 - val_loss: 0.3302 - val_acc: 0.8877 -
val_auc: 0.9830 - val_f1_score: 0.8886
Epoch 23/100
256/256 [=====] - 13s 51ms/step - loss: 0.1199 - acc:
0.9572 - auc: 0.9966 - f1_score: 0.9571 - val_loss: 0.8247 - val_acc: 0.7896 -
val_auc: 0.9399 - val_f1_score: 0.7841
Epoch 24/100
256/256 [=====] - 13s 51ms/step - loss: 0.1231 - acc:
0.9586 - auc: 0.9961 - f1_score: 0.9584 - val_loss: 0.2141 - val_acc: 0.9233 -
val_auc: 0.9920 - val_f1_score: 0.9231
Epoch 25/100
256/256 [=====] - 13s 51ms/step - loss: 0.1011 - acc:

0.9662 - auc: 0.9977 - f1_score: 0.9660 - val_loss: 0.6412 - val_acc: 0.8159 - val_auc: 0.9583 - val_f1_score: 0.8014

Epoch 26/100

256/256 [=====] - 13s 51ms/step - loss: 0.1094 - acc: 0.9643 - auc: 0.9971 - f1_score: 0.9643 - val_loss: 0.2493 - val_acc: 0.9170 - val_auc: 0.9902 - val_f1_score: 0.9160

Epoch 27/100

256/256 [=====] - 13s 51ms/step - loss: 0.0896 - acc: 0.9669 - auc: 0.9982 - f1_score: 0.9668 - val_loss: 0.5295 - val_acc: 0.8286 - val_auc: 0.9687 - val_f1_score: 0.8225

Epoch 28/100

256/256 [=====] - 13s 51ms/step - loss: 0.1001 - acc: 0.9646 - auc: 0.9972 - f1_score: 0.9648 - val_loss: 0.3778 - val_acc: 0.8853 - val_auc: 0.9805 - val_f1_score: 0.8820

Epoch 29/100

256/256 [=====] - 13s 51ms/step - loss: 0.0986 - acc: 0.9687 - auc: 0.9971 - f1_score: 0.9685 - val_loss: 0.3089 - val_acc: 0.9033 - val_auc: 0.9845 - val_f1_score: 0.9013

Epoch 30/100

256/256 [=====] - 13s 51ms/step - loss: 0.0767 - acc: 0.9705 - auc: 0.9987 - f1_score: 0.9708 - val_loss: 0.5289 - val_acc: 0.8496 - val_auc: 0.9686 - val_f1_score: 0.8403

Epoch 31/100

256/256 [=====] - 13s 51ms/step - loss: 0.0695 - acc: 0.9790 - auc: 0.9982 - f1_score: 0.9789 - val_loss: 0.2920 - val_acc: 0.9136 - val_auc: 0.9864 - val_f1_score: 0.9119

Epoch 32/100

256/256 [=====] - 13s 51ms/step - loss: 0.0705 - acc: 0.9768 - auc: 0.9985 - f1_score: 0.9766 - val_loss: 0.3391 - val_acc: 0.9033 - val_auc: 0.9828 - val_f1_score: 0.8998

Epoch 33/100

256/256 [=====] - 13s 51ms/step - loss: 0.0656 - acc: 0.9800 - auc: 0.9985 - f1_score: 0.9799 - val_loss: 1.2103 - val_acc: 0.7383 - val_auc: 0.9036 - val_f1_score: 0.7300

Epoch 34/100

256/256 [=====] - 13s 51ms/step - loss: 0.0620 - acc: 0.9798 - auc: 0.9985 - f1_score: 0.9799 - val_loss: 0.7214 - val_acc: 0.8228 - val_auc: 0.9495 - val_f1_score: 0.8189

Epoch 35/100

256/256 [=====] - 13s 51ms/step - loss: 0.0696 - acc: 0.9765 - auc: 0.9987 - f1_score: 0.9764 - val_loss: 0.4714 - val_acc: 0.8784 - val_auc: 0.9711 - val_f1_score: 0.8720

Epoch 36/100

256/256 [=====] - 13s 51ms/step - loss: 0.0575 - acc: 0.9814 - auc: 0.9986 - f1_score: 0.9814 - val_loss: 0.7210 - val_acc: 0.8198 - val_auc: 0.9489 - val_f1_score: 0.8091

Epoch 37/100

256/256 [=====] - 13s 51ms/step - loss: 0.0631 - acc:

0.9804 - auc: 0.9987 - f1_score: 0.9802 - val_loss: 0.2202 - val_acc: 0.9355 -
val_auc: 0.9900 - val_f1_score: 0.9351
Epoch 38/100
256/256 [=====] - 13s 51ms/step - loss: 0.0979 - acc:
0.9676 - auc: 0.9967 - f1_score: 0.9677 - val_loss: 0.2414 - val_acc: 0.9136 -
val_auc: 0.9898 - val_f1_score: 0.9123
Epoch 39/100
256/256 [=====] - 13s 51ms/step - loss: 0.0725 - acc:
0.9763 - auc: 0.9986 - f1_score: 0.9763 - val_loss: 0.2160 - val_acc: 0.9375 -
val_auc: 0.9886 - val_f1_score: 0.9375
Epoch 40/100
256/256 [=====] - 13s 51ms/step - loss: 0.0511 - acc:
0.9817 - auc: 0.9993 - f1_score: 0.9815 - val_loss: 0.2085 - val_acc: 0.9346 -
val_auc: 0.9907 - val_f1_score: 0.9347
Epoch 41/100
256/256 [=====] - 13s 51ms/step - loss: 0.0359 - acc:
0.9877 - auc: 0.9994 - f1_score: 0.9877 - val_loss: 0.5973 - val_acc: 0.8540 -
val_auc: 0.9605 - val_f1_score: 0.8555
Epoch 42/100
256/256 [=====] - 13s 51ms/step - loss: 0.0381 - acc:
0.9880 - auc: 0.9994 - f1_score: 0.9879 - val_loss: 0.2313 - val_acc: 0.9365 -
val_auc: 0.9895 - val_f1_score: 0.9367
Epoch 43/100
256/256 [=====] - 13s 51ms/step - loss: 0.0378 - acc:
0.9877 - auc: 0.9996 - f1_score: 0.9876 - val_loss: 0.4954 - val_acc: 0.8794 -
val_auc: 0.9677 - val_f1_score: 0.8761
Epoch 44/100
256/256 [=====] - 13s 51ms/step - loss: 0.0646 - acc:
0.9782 - auc: 0.9986 - f1_score: 0.9782 - val_loss: 0.2919 - val_acc: 0.9219 -
val_auc: 0.9850 - val_f1_score: 0.9207
Epoch 45/100
256/256 [=====] - 13s 51ms/step - loss: 0.0505 - acc:
0.9855 - auc: 0.9991 - f1_score: 0.9855 - val_loss: 0.2164 - val_acc: 0.9395 -
val_auc: 0.9895 - val_f1_score: 0.9395
Epoch 46/100
256/256 [=====] - 13s 51ms/step - loss: 0.0284 - acc:
0.9915 - auc: 0.9995 - f1_score: 0.9914 - val_loss: 0.5887 - val_acc: 0.8472 -
val_auc: 0.9596 - val_f1_score: 0.8423
Epoch 47/100
256/256 [=====] - 13s 51ms/step - loss: 0.0243 - acc:
0.9927 - auc: 0.9997 - f1_score: 0.9927 - val_loss: 0.3140 - val_acc: 0.9067 -
val_auc: 0.9842 - val_f1_score: 0.9048
Epoch 48/100
256/256 [=====] - 13s 51ms/step - loss: 0.0662 - acc:
0.9809 - auc: 0.9980 - f1_score: 0.9809 - val_loss: 0.2009 - val_acc: 0.9458 -
val_auc: 0.9890 - val_f1_score: 0.9458
Epoch 49/100
256/256 [=====] - 13s 51ms/step - loss: 0.0363 - acc:

0.9877 - auc: 0.9997 - f1_score: 0.9877 - val_loss: 0.2108 - val_acc: 0.9375 -
val_auc: 0.9904 - val_f1_score: 0.9377

Epoch 50/100

256/256 [=====] - 13s 51ms/step - loss: 0.0276 - acc:
0.9918 - auc: 0.9993 - f1_score: 0.9918 - val_loss: 0.9958 - val_acc: 0.7876 -
val_auc: 0.9227 - val_f1_score: 0.7801

Epoch 51/100

256/256 [=====] - 13s 51ms/step - loss: 0.0376 - acc:
0.9888 - auc: 0.9996 - f1_score: 0.9888 - val_loss: 0.2587 - val_acc: 0.9297 -
val_auc: 0.9867 - val_f1_score: 0.9287

Epoch 52/100

256/256 [=====] - 13s 51ms/step - loss: 0.0343 - acc:
0.9892 - auc: 0.9993 - f1_score: 0.9893 - val_loss: 0.7595 - val_acc: 0.8032 -
val_auc: 0.9456 - val_f1_score: 0.7983

Epoch 53/100

256/256 [=====] - 13s 51ms/step - loss: 0.0711 - acc:
0.9789 - auc: 0.9980 - f1_score: 0.9788 - val_loss: 0.2659 - val_acc: 0.9253 -
val_auc: 0.9849 - val_f1_score: 0.9245

Epoch 54/100

256/256 [=====] - 13s 51ms/step - loss: 0.0433 - acc:
0.9858 - auc: 0.9995 - f1_score: 0.9859 - val_loss: 2.8259 - val_acc: 0.5425 -
val_auc: 0.7920 - val_f1_score: 0.4884

Epoch 55/100

256/256 [=====] - 13s 51ms/step - loss: 0.0412 - acc:
0.9882 - auc: 0.9993 - f1_score: 0.9883 - val_loss: 0.2718 - val_acc: 0.9307 -
val_auc: 0.9849 - val_f1_score: 0.9303

Epoch 56/100

256/256 [=====] - 13s 51ms/step - loss: 0.0406 - acc:
0.9875 - auc: 0.9987 - f1_score: 0.9875 - val_loss: 0.1802 - val_acc: 0.9482 -
val_auc: 0.9901 - val_f1_score: 0.9479

Epoch 57/100

256/256 [=====] - 13s 51ms/step - loss: 0.0243 - acc:
0.9920 - auc: 0.9996 - f1_score: 0.9920 - val_loss: 0.5178 - val_acc: 0.8696 -
val_auc: 0.9680 - val_f1_score: 0.8671

Epoch 58/100

256/256 [=====] - 13s 51ms/step - loss: 0.0389 - acc:
0.9879 - auc: 0.9989 - f1_score: 0.9880 - val_loss: 0.4408 - val_acc: 0.8872 -
val_auc: 0.9758 - val_f1_score: 0.8839

Epoch 59/100

256/256 [=====] - 13s 51ms/step - loss: 0.0432 - acc:
0.9858 - auc: 0.9992 - f1_score: 0.9858 - val_loss: 0.5958 - val_acc: 0.8691 -
val_auc: 0.9607 - val_f1_score: 0.8609

Epoch 60/100

256/256 [=====] - 13s 51ms/step - loss: 0.0356 - acc:
0.9865 - auc: 0.9997 - f1_score: 0.9864 - val_loss: 0.4801 - val_acc: 0.8887 -
val_auc: 0.9693 - val_f1_score: 0.8860

Epoch 61/100

256/256 [=====] - 13s 51ms/step - loss: 0.0286 - acc:

0.9883 - auc: 0.9998 - f1_score: 0.9884 - val_loss: 0.2878 - val_acc: 0.9365 -
val_auc: 0.9827 - val_f1_score: 0.9364

Epoch 62/100

256/256 [=====] - 13s 51ms/step - loss: 0.0309 - acc:
0.9900 - auc: 0.9993 - f1_score: 0.9899 - val_loss: 0.2086 - val_acc: 0.9380 -
val_auc: 0.9908 - val_f1_score: 0.9373

Epoch 63/100

256/256 [=====] - 13s 51ms/step - loss: 0.0283 - acc:
0.9914 - auc: 0.9996 - f1_score: 0.9914 - val_loss: 0.3530 - val_acc: 0.9175 -
val_auc: 0.9787 - val_f1_score: 0.9157

Epoch 64/100

256/256 [=====] - 13s 51ms/step - loss: 0.0310 - acc:
0.9906 - auc: 0.9995 - f1_score: 0.9908 - val_loss: 0.2598 - val_acc: 0.9302 -
val_auc: 0.9873 - val_f1_score: 0.9300

Epoch 65/100

256/256 [=====] - 13s 51ms/step - loss: 0.0160 - acc:
0.9951 - auc: 0.9999 - f1_score: 0.9951 - val_loss: 0.2890 - val_acc: 0.9341 -
val_auc: 0.9827 - val_f1_score: 0.9331

Epoch 66/100

256/256 [=====] - 13s 51ms/step - loss: 0.0298 - acc:
0.9901 - auc: 0.9997 - f1_score: 0.9901 - val_loss: 0.3373 - val_acc: 0.9155 -
val_auc: 0.9793 - val_f1_score: 0.9138

Epoch 67/100

256/256 [=====] - 13s 51ms/step - loss: 0.0288 - acc:
0.9928 - auc: 0.9994 - f1_score: 0.9928 - val_loss: 0.3290 - val_acc: 0.9126 -
val_auc: 0.9827 - val_f1_score: 0.9128

Epoch 68/100

256/256 [=====] - 13s 51ms/step - loss: 0.0319 - acc:
0.9899 - auc: 0.9992 - f1_score: 0.9899 - val_loss: 0.1925 - val_acc: 0.9453 -
val_auc: 0.9900 - val_f1_score: 0.9451

Epoch 69/100

256/256 [=====] - 13s 51ms/step - loss: 0.0307 - acc:
0.9897 - auc: 0.9994 - f1_score: 0.9896 - val_loss: 0.2590 - val_acc: 0.9395 -
val_auc: 0.9855 - val_f1_score: 0.9386

Epoch 70/100

256/256 [=====] - 13s 51ms/step - loss: 0.0274 - acc:
0.9915 - auc: 0.9998 - f1_score: 0.9916 - val_loss: 0.2854 - val_acc: 0.9297 -
val_auc: 0.9841 - val_f1_score: 0.9298

Epoch 71/100

256/256 [=====] - 13s 51ms/step - loss: 0.0233 - acc:
0.9929 - auc: 0.9998 - f1_score: 0.9928 - val_loss: 0.2493 - val_acc: 0.9365 -
val_auc: 0.9872 - val_f1_score: 0.9363

Epoch 72/100

256/256 [=====] - 13s 51ms/step - loss: 0.0320 - acc:
0.9921 - auc: 0.9991 - f1_score: 0.9921 - val_loss: 0.2800 - val_acc: 0.9282 -
val_auc: 0.9846 - val_f1_score: 0.9270

Epoch 73/100

256/256 [=====] - 13s 51ms/step - loss: 0.0157 - acc:

0.9950 - auc: 0.9998 - f1_score: 0.9951 - val_loss: 0.2947 - val_acc: 0.9307 -
val_auc: 0.9817 - val_f1_score: 0.9309

Epoch 74/100

256/256 [=====] - 13s 51ms/step - loss: 0.0263 - acc:
0.9896 - auc: 0.9995 - f1_score: 0.9893 - val_loss: 0.3077 - val_acc: 0.9175 -
val_auc: 0.9836 - val_f1_score: 0.9160

Epoch 75/100

256/256 [=====] - 13s 51ms/step - loss: 0.0262 - acc:
0.9919 - auc: 0.9997 - f1_score: 0.9919 - val_loss: 0.8139 - val_acc: 0.8296 -
val_auc: 0.9400 - val_f1_score: 0.8270

Epoch 76/100

256/256 [=====] - 13s 51ms/step - loss: 0.0192 - acc:
0.9925 - auc: 0.9998 - f1_score: 0.9925 - val_loss: 0.2215 - val_acc: 0.9385 -
val_auc: 0.9912 - val_f1_score: 0.9380

Epoch 77/100

256/256 [=====] - 13s 51ms/step - loss: 0.0422 - acc:
0.9855 - auc: 0.9994 - f1_score: 0.9856 - val_loss: 0.2192 - val_acc: 0.9429 -
val_auc: 0.9885 - val_f1_score: 0.9424

Epoch 78/100

256/256 [=====] - 13s 51ms/step - loss: 0.0215 - acc:
0.9942 - auc: 0.9996 - f1_score: 0.9941 - val_loss: 0.1965 - val_acc: 0.9546 -
val_auc: 0.9870 - val_f1_score: 0.9545

Epoch 79/100

256/256 [=====] - 13s 51ms/step - loss: 0.0124 - acc:
0.9960 - auc: 0.9998 - f1_score: 0.9960 - val_loss: 0.2881 - val_acc: 0.9341 -
val_auc: 0.9841 - val_f1_score: 0.9329

Epoch 80/100

256/256 [=====] - 13s 51ms/step - loss: 0.0209 - acc:
0.9926 - auc: 0.9997 - f1_score: 0.9926 - val_loss: 0.2738 - val_acc: 0.9277 -
val_auc: 0.9863 - val_f1_score: 0.9268

Epoch 81/100

256/256 [=====] - 13s 51ms/step - loss: 0.0163 - acc:
0.9952 - auc: 0.9998 - f1_score: 0.9951 - val_loss: 0.2016 - val_acc: 0.9487 -
val_auc: 0.9896 - val_f1_score: 0.9484

Epoch 82/100

256/256 [=====] - 13s 51ms/step - loss: 0.0168 - acc:
0.9947 - auc: 0.9998 - f1_score: 0.9947 - val_loss: 0.5604 - val_acc: 0.8687 -
val_auc: 0.9623 - val_f1_score: 0.8674

Epoch 83/100

256/256 [=====] - 13s 51ms/step - loss: 0.0282 - acc:
0.9906 - auc: 0.9994 - f1_score: 0.9906 - val_loss: 0.3112 - val_acc: 0.9312 -
val_auc: 0.9812 - val_f1_score: 0.9299

Epoch 84/100

256/256 [=====] - 13s 51ms/step - loss: 0.0322 - acc:
0.9917 - auc: 0.9990 - f1_score: 0.9917 - val_loss: 0.7490 - val_acc: 0.8315 -
val_auc: 0.9482 - val_f1_score: 0.8284

Epoch 85/100

256/256 [=====] - 13s 51ms/step - loss: 0.0188 - acc:

0.9945 - auc: 0.9997 - f1_score: 0.9945 - val_loss: 0.2061 - val_acc: 0.9492 -
val_auc: 0.9887 - val_f1_score: 0.9493

Epoch 86/100

256/256 [=====] - 13s 51ms/step - loss: 0.0163 - acc:
0.9950 - auc: 0.9998 - f1_score: 0.9950 - val_loss: 0.2150 - val_acc: 0.9443 -
val_auc: 0.9896 - val_f1_score: 0.9438

Epoch 87/100

256/256 [=====] - 13s 51ms/step - loss: 0.0165 - acc:
0.9943 - auc: 0.9998 - f1_score: 0.9940 - val_loss: 0.2274 - val_acc: 0.9414 -
val_auc: 0.9875 - val_f1_score: 0.9411

Epoch 88/100

256/256 [=====] - 13s 51ms/step - loss: 0.0146 - acc:
0.9956 - auc: 0.9999 - f1_score: 0.9956 - val_loss: 0.4331 - val_acc: 0.9067 -
val_auc: 0.9733 - val_f1_score: 0.9058

Epoch 89/100

256/256 [=====] - 13s 51ms/step - loss: 0.0125 - acc:
0.9958 - auc: 0.9999 - f1_score: 0.9959 - val_loss: 0.2582 - val_acc: 0.9409 -
val_auc: 0.9850 - val_f1_score: 0.9413

Epoch 90/100

256/256 [=====] - 13s 51ms/step - loss: 0.0201 - acc:
0.9944 - auc: 0.9998 - f1_score: 0.9944 - val_loss: 0.2714 - val_acc: 0.9434 -
val_auc: 0.9831 - val_f1_score: 0.9425

Epoch 91/100

256/256 [=====] - 13s 51ms/step - loss: 0.0087 - acc:
0.9973 - auc: 0.9999 - f1_score: 0.9973 - val_loss: 0.4787 - val_acc: 0.9072 -
val_auc: 0.9672 - val_f1_score: 0.9059

Epoch 92/100

256/256 [=====] - 13s 51ms/step - loss: 0.0219 - acc:
0.9924 - auc: 0.9998 - f1_score: 0.9923 - val_loss: 0.3455 - val_acc: 0.9219 -
val_auc: 0.9780 - val_f1_score: 0.9194

Epoch 93/100

256/256 [=====] - 13s 51ms/step - loss: 0.0209 - acc:
0.9927 - auc: 0.9999 - f1_score: 0.9927 - val_loss: 0.2662 - val_acc: 0.9365 -
val_auc: 0.9844 - val_f1_score: 0.9361

Epoch 94/100

256/256 [=====] - 13s 51ms/step - loss: 0.0228 - acc:
0.9933 - auc: 0.9995 - f1_score: 0.9933 - val_loss: 0.5512 - val_acc: 0.8857 -
val_auc: 0.9640 - val_f1_score: 0.8822

Epoch 95/100

256/256 [=====] - 13s 51ms/step - loss: 0.0133 - acc:
0.9969 - auc: 0.9996 - f1_score: 0.9969 - val_loss: 0.2214 - val_acc: 0.9502 -
val_auc: 0.9879 - val_f1_score: 0.9501

Epoch 96/100

256/256 [=====] - 13s 51ms/step - loss: 0.0151 - acc:
0.9948 - auc: 1.0000 - f1_score: 0.9948 - val_loss: 0.3272 - val_acc: 0.9248 -
val_auc: 0.9795 - val_f1_score: 0.9238

Epoch 97/100

256/256 [=====] - 13s 51ms/step - loss: 0.0110 - acc:

```

0.9957 - auc: 0.9999 - f1_score: 0.9956 - val_loss: 0.2312 - val_acc: 0.9429 -
val_auc: 0.9875 - val_f1_score: 0.9431
Epoch 98/100
256/256 [=====] - 13s 51ms/step - loss: 0.0132 - acc:
0.9955 - auc: 0.9996 - f1_score: 0.9955 - val_loss: 0.2145 - val_acc: 0.9502 -
val_auc: 0.9875 - val_f1_score: 0.9501
Epoch 99/100
256/256 [=====] - 13s 51ms/step - loss: 0.0092 - acc:
0.9972 - auc: 1.0000 - f1_score: 0.9972 - val_loss: 0.4792 - val_acc: 0.9019 -
val_auc: 0.9705 - val_f1_score: 0.9004
Epoch 100/100
256/256 [=====] - 13s 51ms/step - loss: 0.0167 - acc:
0.9952 - auc: 0.9998 - f1_score: 0.9952 - val_loss: 0.2442 - val_acc: 0.9482 -
val_auc: 0.9849 - val_f1_score: 0.9481

```


[16]: *#Plotting the trend of the metrics during training*

```

fig, ax = plt.subplots(1, 3, figsize = (30, 5))
ax = ax.ravel()

for i, metric in enumerate(["acc", "auc", "loss"]):
    ax[i].plot(history.history[metric])
    ax[i].plot(history.history["val_" + metric])
    ax[i].set_title("Model {}".format(metric))
    ax[i].set_xlabel("Epochs")
    ax[i].set_ylabel(metric)
    ax[i].legend(["train", "val"])

```



[17]: *#Evaluating the model on the data*

```

#train_scores = model.evaluate(train_data, train_labels)
#val_scores = model.evaluate(val_data, val_labels)
test_scores = model.evaluate(test_data, test_labels)

#print("Training Accuracy: %.2f%%"%(train_scores[1] * 100))
#print("Validation Accuracy: %.2f%%"%(val_scores[1] * 100))
print("Testing Accuracy: %.2f%%"%(test_scores[1] * 100))

```

```

80/80 [=====] - 1s 14ms/step - loss: 0.2458 - acc:

```


0.9488 - auc: 0.9850 - f1_score: 0.9493
Testing Accuracy: 94.88%

```
[18]: #Predicting the test data
```

```
pred_labels = model.predict(test_data)
```

```
[19]: #Print the classification report of the tested data
```

```
#Since the labels are softmax arrays, we need to roundoff to have it in the
↳ form of 0s and 1s,
#similar to the test_labels
def roundoff(arr):
    """To round off according to the argmax of each predicted label array. """
    arr[np.argmax(arr) != arr.max()] = 0
    arr[np.argmax(arr) == arr.max()] = 1
    return arr

for labels in pred_labels:
    labels = roundoff(labels)

print(classification_report(test_labels, pred_labels, target_names=CLASSES))
```

	precision	recall	f1-score	support
NonDemented	0.99	0.95	0.97	639
VeryMildDemented	1.00	1.00	1.00	635
MildDemented	0.94	0.89	0.91	662
ModerateDemented	0.88	0.95	0.91	624
micro avg	0.95	0.95	0.95	2560
macro avg	0.95	0.95	0.95	2560
weighted avg	0.95	0.95	0.95	2560
samples avg	0.95	0.95	0.95	2560

```
[20]: #Plot the confusion matrix to understand the classification in detail
```


```
pred_ls = np.argmax(pred_labels, axis=1)
test_ls = np.argmax(test_labels, axis=1)

conf_arr = confusion_matrix(test_ls, pred_ls)

plt.figure(figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

ax = sns.heatmap(conf_arr, cmap='Greens', annot=True, fmt='d',
↳ xticklabels=CLASSES, yticklabels=CLASSES)
```

```
plt.title('Alzheimer\'s Disease Diagnosis')
plt.xlabel('Prediction')
plt.ylabel('Truth')
plt.show(ax)
```



[21]: *#Printing some other classification metrics*

```
print("Balanced Accuracy Score: {} %".format(round(BAS(test_ls, pred_ls) * 100, 2)))
print("Matthew's Correlation Coefficient: {} %".format(round(MCC(test_ls, pred_ls) * 100, 2)))
```

Balanced Accuracy Score: 94.94 %

Matthew's Correlation Coefficient: 93.22 %

[22]: *#Saving the model for future use*

```
model_dir = work_dir + "alzheimer_cnn_model"
model.save(model_dir, save_format='h5')
os.listdir(work_dir)
```

```
[22]: ['alzheimer_cnn_model',  
      'MildDemented',  
      'NonDemented',  
      'ModerateDemented',  
      'VeryMildDemented']
```

```
[23]: pretrained_model = tf.keras.models.load_model(model_dir)  
  
      #Check its architecture  
      plot_model(pretrained_model, to_file=work_dir + "model_plot.png",  
                  ↪show_shapes=True, show_layer_names=True)
```

```
[23]:
```

